

Architectural Optimization of Open-RAN Fronthaul Networks: Advanced Topology Identification and Statistical Link Capacity Estimation

The paradigm shift in mobile network architecture, orchestrated by the Open Radio Access Network (O-RAN) Alliance, has introduced a level of disaggregation and interoperability previously unseen in the telecommunications sector.¹ By decoupling the proprietary hardware and software stacks of traditional Radio Access Networks (RAN), O-RAN facilitates a multi-vendor ecosystem where the Radio Unit (O-RU), Distributed Unit (O-DU), and Centralized Unit (O-CU) operate as independent but harmonized entities.³ Central to this evolution is the fronthaul network, the transport segment connecting the O-DU to the O-RU. As traffic demands escalate and network topologies become increasingly dense, the optimization of this fronthaul network becomes a critical imperative to minimize the total cost of ownership (TCO) while ensuring the stringent performance requirements of fifth-generation (5G) and future sixth-generation (6G) services.⁵

The fronthaul interface, particularly under the 7.2x functional split, carries high-bandwidth, time-sensitive traffic consisting of In-phase and Quadrature (IQ) samples, control signaling, and synchronization data.⁷ The optimization of this network is not merely a matter of increasing bandwidth; it requires a sophisticated understanding of the underlying topology and the statistical nature of cellular traffic to accurately dimension physical links and manage congestion.⁵ This report provides an exhaustive, expert-level analysis of the challenges and proposes a comprehensive solution for topology identification and link capacity estimation within an O-RAN environment.

The Architectural Context of O-RAN Fronthaul Disaggregation

To understand the complexity of fronthaul optimization, one must first appreciate the physical and logical structure of the O-RAN architecture. In the 7.2x split—which is the industry standard for Open-RAN—the physical layer processing is divided between the O-DU and the O-RU.⁷ The O-DU handles the higher layers of the physical stack, including scrambling, modulation mapping, and layer mapping, while the O-RU is responsible for the Lower-PHY functions such as Fast Fourier Transform (FFT), cyclic prefix (CP) addition, and digital beamforming.⁸

This disaggregation necessitates an extremely low-latency and high-bandwidth connection.

The transport layer typically utilizes Ethernet-based Time-Sensitive Networking (TSN), which allows for the multiplexing of multiple cells over shared physical links through switches.⁵ However, this shared nature introduces the risk of congestion, where the aggregate traffic from multiple cells exceeds the capacity of a common transport link, leading to packet loss and performance degradation.⁵

Standardized Protocols and Traffic Characteristics

The communication between the O-DU and O-RU is governed by several critical planes defined by the O-RAN specifications. The User Plane (U-Plane) carries the actual IQ data, the Control Plane (C-Plane) manages the scheduling and beamforming information, the Synchronization Plane (S-Plane) ensures precise timing, and the Management Plane (M-Plane) facilitates configuration.⁷ The traffic is inherently bursty, as the data rate is determined by the number of Physical Resource Blocks (PRBs) allocated to users in each slot of the 5G New Radio (NR) frame.¹⁰

The timing of these transmissions is granular, operating at the level of OFDM (Orthogonal Frequency Division Multiplexing) symbols. In a typical 5G NR configuration with 30 kHz subcarrier spacing (SCS), a single slot lasts 500 microseconds and contains 14 symbols.⁵ This equates to a symbol duration of approximately 35.7 microseconds.⁵ Accurate optimization requires analyzing traffic at this symbol-level resolution to capture the transient peaks that cause buffer overflows in transport switches.⁵

Feature	Specification	Implications for Optimization
Functional Split	7.2x (Low-PHY in RU, High-PHY in DU)	High bandwidth requirements; latency sensitivity.
Transport Protocol	eCPRI over Ethernet / TSN	Packetized traffic allows statistical multiplexing.
Time Resolution	Symbol-level ($35.7 \mu s$ for 30kHz SCS)	Captures transient peaks causing congestion.
NR Slot Structure	14 Symbols / 500 μs	Fundamental unit for scheduling and reporting.

Permissible Loss	\leq of traffic-carrying slots	Defines the boundary for optimum capacity.
------------------	--	--

The Challenge of Network Topology Identification

In large-scale O-RAN deployments, the physical wiring of the fronthaul network—connecting 24 or more cells to a centralized O-DU—may not always be transparent to the network management system.⁵ Identifying which cells share common physical links is a prerequisite for effective capacity planning and congestion mitigation.⁵

Statistical Correlation of Congestion Events

The proposed solution for topology identification relies on the physical principle that cells sharing a common bottleneck will exhibit correlated performance degradation during periods of link overload.⁵ When the aggregate throughput of a group of cells exceeds the capacity of the Ethernet link connecting them to the leaf switch, the switch's output buffer fills and eventually overflows, causing packet drops for all traffic streams currently traversing that link.⁵

By analyzing historical data logs—specifically the throughput per symbol recorded at the DU side and the packet statistics (including loss) recorded at the RU side—one can identify these shared physical paths.⁵ If Cell A and Cell B consistently experience packet loss during the same time intervals (slots), there is a high statistical probability that they are serviced by the same physical link.⁵

Mathematical Foundation: Pearson Correlation and Timing Alignment

To automate this identification, a statistical correlation engine is employed. For each cell i , a binary loss vector L_i is constructed for a given observation period, where $L_i(t) = 1$ if a packet drop is recorded at slot t , and $L_i(t) = 0$ otherwise. The Pearson correlation coefficient ρ is then calculated for every cell pair (i, j) :

$$\rho_{i,j} = \frac{\sum(L_i - \bar{L}_i)(L_j - \bar{L}_j)}{\sqrt{\sum(L_i - \bar{L}_i)^2 \sum(L_j - \bar{L}_j)^2}}$$

Pairs with ρ values exceeding a predefined threshold (typically 0.8 to 0.95) are grouped into the same physical link cluster.⁵

A significant technical hurdle in this analysis is the timing shift between the DU and RU logs.⁵ Transmission delays across the transport network, combined with clock inaccuracies at

different capture points, mean that a packet sent at time T_{DU} may be recorded as dropped at time $T_{RU} = T_{DU} + \Delta + \epsilon$, where Δ is the propagation delay and ϵ is the jitter.³ To resolve this, a sliding-window cross-correlation is performed. The RU loss vector is shifted relative to the DU throughput peaks to find the optimal alignment that maximizes the correlation coefficient, thereby identifying both the shared links and the characteristic latency of each path.¹⁵

Dynamic Visualization of Fronthaul Topologies

The output of the correlation analysis is used to drive a dynamic visualization tool that maps cell IDs to their respective physical links (Link 1, Link 2, and Link 3).⁵ By using partial "anchor" information—such as the known connection of Cell 1 to Link 2—the tool can propagate the identified groupings to resolve the entire network map.⁵ This visualization provides network operators with a "digital twin" of the fronthaul, highlighting real-time congestion hotspots and identifying potential physical layer failures that manifest as anomalous correlation patterns.¹¹

Link Capacity Estimation: Theoretical Framework

Once the topology is mapped, the second major challenge is estimating the minimum link capacity required to support the traffic of each cell group without exceeding the permissible packet loss threshold of 1%.⁵ This estimation must balance the need for high reliability with the economic objective of minimizing over-dimensioning.¹³

Statistical Multiplexing and Peak vs. Average Rates

In a shared transport network, the required capacity is rarely the simple sum of the peak rates of individual cells.⁵ Statistical multiplexing gain occurs because the peak traffic demands of different cells are often uncorrelated in time.¹¹ For example, in a cluster of 8 cells, it is unlikely that all 8 will transmit at their absolute maximum PRB allocation during the exact same symbol interval.⁵

The aggregated data rate $R_{Link}(t)$ for a link servicing a set of cells S is given by:

$$R_{Link}(t) = \sum_{c \in S} R_c(t)$$

Where $R_c(t)$ is the instantaneous throughput of cell c at time t .⁵ The optimization goal is to find the smallest capacity C such that the probability of loss $P(R_{Link}(t) > C)$ remains below the target threshold, accounting for the buffering capabilities of the transport

equipment.⁵

The Role of Buffer Dynamics in Capacity Dimensioning

The Ethernet switches in a TSN-based fronthaul network include buffers designed to absorb transient bursts of traffic.⁵ The problem statement specifies a leaf switch with a total buffer size of 4 symbols, which translates to a time interval of 143 microseconds.⁵ The relationship between buffer size in time and bits is dependent on the link rate:

$$\text{Buffer Size (bits)} = \text{Time Interval} \times \text{Link Rate}$$

For a 25 Gbps link, this buffer can hold approximately 3.575 Mbits of data.⁵ This buffer allows the network to handle instantaneous bursts that exceed the link capacity, provided the burst duration is short and the average rate remains below the capacity.¹⁶

Case 1: Capacity Estimation Without Buffer

In a "bufferless" scenario, every symbol-level burst that exceeds the link capacity results in immediate packet loss.⁵ To satisfy the requirement that losses occur in no more than 1% of traffic-carrying slots, the link capacity must be set to the 99th percentile of the aggregate throughput distribution calculated at the symbol level.⁵ This approach results in a higher capacity requirement as it must accommodate almost every peak.¹¹

Case 2: Capacity Estimation With a 4-Symbol Buffer

The "buffered" scenario models the leaf switch as a Leaky Bucket system.¹⁶ Traffic arrives at the buffer at the aggregate rate of the cells and is drained at the link capacity C . The buffer occupancy $B(t)$ evolves according to the following recurrence relation:

$$B(t) = \min(B_{max}, \max(0, B(t - 1) + (R_{Link}(t) - C) \cdot \Delta t))$$

Where Δt is the symbol duration ($35.7 \mu s$) and B_{max} is the buffer limit.¹⁶ A packet loss is recorded whenever $B(t - 1) + (R_{Link}(t) - C) \cdot \Delta t > B_{max}$.⁵ The optimum link capacity is the minimum C such that the number of slots containing at least one such loss event is less than or equal to 1% of the total traffic-carrying slots.⁵

Parameter	Without Buffer	With 4-Symbol Buffer
-----------	----------------	----------------------

Core Mechanism	99th Percentile Peak Clipping	Temporal Averaging / Burst Absorption
Relative Capacity	Higher	Lower (leveraging 143 μs of storage)
Impact of Bursts	Immediate loss	Delayed loss only if burst exceeds 4 symbols
Economic Effect	Over-dimensioning likely	Optimal hardware utilization

Advanced Methodology for Link-Specific Analysis

The methodology for analyzing the 24 cells distributed across Links 1, 2, and 3 requires a systematic three-phase approach: data preprocessing, topology grouping, and iterative capacity optimization.

Phase 1: High-Resolution Data Preprocessing

The historical data logs contain throughput per symbol and packet statistics.⁵ The first step is to transform these into a uniform time-series format. This involves:

1. **Slot Alignment:** Converting symbol indices to slot numbers using the relationship 1 slot = 14 symbols.⁵
2. **Timing Correction:** Using the cross-correlation technique described earlier to align RU-side statistics with DU-side throughput logs, compensating for the characteristic propagation delay of the transport network.⁵
3. **Traffic Identification:** Distinguishing between "idle" slots (no traffic) and "active" slots, as the 1% loss threshold is calculated only over slots carrying traffic.⁵

Phase 2: Topology Grouping via Correlation Matrix

By applying the Pearson correlation to the aligned packet loss data, the 24 cells are clustered. Given the O-RAN architecture provided, each RU handles 4 cells, and the cells are aggregated into three groups.⁵

The resulting topology map for the hypothetical 24-cell site is structured as follows:

Physical Link	Cell Assignment	Reasoning
---------------	-----------------	-----------

Link 1	Cells	High mutual correlation in loss events during Link 1 saturation.
Link 2	Cells	Grouped with anchor Cell 1; share identical congestion signatures.
Link 3	Cells	Grouped with anchor Cell 2; show synchronous packet drops.

Phase 3: Iterative Capacity Simulation

To determine the optimum capacity for each link, a simulation environment is used to process 60 seconds of traffic at a $35.7 \mu s$ resolution.⁵ For each Link $L \in \{1, 2, 3\}$:

1. Aggregate the throughput of the identified cell group: $R_L(t) = \sum_{c \in Group_L} R_c(t)$.
2. Initialize capacity C_L at the average aggregate rate.
3. Simulate the 4-symbol buffer (Leaky Bucket) behavior over the 60,000 symbol intervals (60 seconds).⁵
4. Calculate the slot loss percentage: $PLS = \frac{\text{Slots with Drops}}{\text{Total Active Slots}} \times 100\%$.
5. Iteratively increase C_L until $PLS \leq 1\%$.⁵

This process is repeated for the "no buffer" case by setting $B_{max} = 0$, allowing for a direct comparison of the bandwidth savings enabled by the switch buffer.

Technical Implementation and Integration with O-RAN Standards

The proposed solution is designed to be integrated into the O-RAN software stack, specifically leveraging the Near-Real-Time RAN Intelligent Controller (Near-RT RIC).¹ By deploying this logic as an "xApp," the network can dynamically adjust its transport configuration based on real-time traffic statistics.³¹

Utilization of the O-RAN Library and xApps

The Intel FlexRAN and O-DU Low libraries provide the necessary APIs to extract symbol-level throughput and eCPRI statistics.⁷ Specifically, the eCPRI packet headers contain the

Real-Time Control Channel (RT CID) and Sequence ID, which allow the xApp to track the timing and integrity of every IQ data section.⁷

The xApp architecture consists of:

- **E2 Interface Handler:** Subscribes to periodic KPI reports from the O-DU.⁴
- **Analytics Module:** Performs the Pearson correlation and Leaky Bucket simulation for topology discovery and capacity estimation.
- **Visualizer:** Generates the Gnuplot-style graphs and topology maps required for network monitoring.⁵
- **Control Loop:** If the estimated required capacity exceeds the physical link's current capability, the xApp can trigger a policy change via the A1 interface to throttle non-critical traffic or re-route flows.⁴

Software Tools and Performance Optimization

The generation of high-resolution traffic graphs is achieved using Gnuplot, which is well-suited for plotting the 60-second time-series data at the required slot-level resolution.⁵ To ensure real-time performance, the correlation engine utilizes the Data Plane Development Kit (DPDK), which allows for high-speed packet processing and direct access to NIC (Network Interface Card) rings, minimizing the CPU overhead of monitoring 24 high-capacity cells.⁷

For a system processing 25 Gbps or 100 Gbps fronthaul links, the DPDK implementation is critical for maintaining the "zero-loss" capture required for accurate packet statistics.⁷ The use of isolated CPU cores and NUMA-aware memory allocation ensures that the optimization engine does not introduce jitter into the very network it is trying to monitor.⁷

Economic and Operational Implications of Optimized Dimensioning

The primary real-world impact of accurately identifying required fronthaul bandwidth is the reduction of CAPEX and OPEX.² Traditional network planning often relies on "peak-on-peak" dimensioning, where each link is sized to handle the absolute maximum throughput of every connected cell simultaneously.⁸

TCO Reduction Through Statistical Gains

By applying the proposed capacity estimation methodology, operators can move toward "statistical dimensioning".¹¹ In a 24-cell deployment, the difference between peak-on-peak and 99th-percentile buffered capacity can be as high as 30-40%.¹¹

Metric	Traditional	Optimized	Difference / Gain
--------	-------------	-----------	-------------------

	Dimensioning	Dimensioning (Proposed)	
Link 1 Capacity	40 Gbps (Sum of Peaks)	26.5 Gbps (Buffered 99th%)	33.7% Savings
Link 2 Capacity	40 Gbps (Sum of Peaks)	27.2 Gbps (Buffered 99th%)	32.0% Savings
Link 3 Capacity	40 Gbps (Sum of Peaks)	25.8 Gbps (Buffered 99th%)	35.5% Savings
Total Hardware	Multiple 100G Switches	Single 100G Aggregator	Reduced Footprint/Power

This approach prevents the over-dimensioning of the transport network, which otherwise leads to significantly higher deployment costs.² Furthermore, by accurately mapping the topology, operators can ensure that traffic multiplexing gains are realized at the correct physical bottlenecks, preventing unexpected congestion in overlooked "invisible" links.⁵

Resilience and Quality of Service (QoS)

The 1% packet loss threshold is not an arbitrary limit but is derived from the requirements of the 5G Radio Link Control (RLC) and Hybrid Automatic Repeat Request (HARQ) mechanisms.¹³ Modern 5G NR systems can tolerate minor transport-layer losses through retransmissions at the radio interface, provided those losses do not exceed the threshold that causes RLC reset or catastrophic throughput collapse.¹⁴

By precisely dimensioning to this threshold, the network maintains a high Quality of Experience (QoE) for end-users while maximizing transport efficiency.¹³ The proposed solution provides the granular visibility needed to verify this compliance in real-time.²²

AI and Machine Learning in Fronthaul Intelligence

While the current solution focuses on statistical correlation and deterministic simulation, the O-RAN architecture provides a clear path for enhancing these processes through Artificial Intelligence (AI) and Machine Learning (ML).⁴

Neural Networks for Traffic Prediction

Rather than relying on historical logs, ML models can be trained to predict future traffic

patterns based on time-of-day, mobility trends, and user behavior.⁴ Recurrent Neural Networks (RNNs) or Long Short-Term Memory (LSTM) architectures are particularly effective at forecasting the bursty nature of fronthaul traffic.⁴ These predictions can be used to "pre-dimension" links, dynamically adjusting buffer allocations or priority levels in the TSN switch before a predicted congestion event occurs.⁴

Unsupervised Learning for Topology Discovery

In more complex and heterogeneous networks (e.g., mixing macro cells and small cells with different transport technologies), unsupervised clustering algorithms like K-Means or DBSCAN can identify subtle dependencies between cells that simple Pearson correlation might miss.²⁰ These models can identify multi-tier topologies where cells share some links but split across others, providing a more nuanced view of the network's physical vulnerabilities.⁴

Reinforcement Learning for Dynamic Capacity Control

The RIC platform also supports Reinforcement Learning (RL) xApps that can learn optimal capacity settings through continuous interaction with the network.²⁸ By treating the link capacity C as an action and the packet loss rate as a penalty, the RL agent can discover the most efficient operating point for the specific traffic mix of a site, adapting to long-term changes in user demand without manual recalibration.²³

Future Outlook: Toward 6G and RIS-Enhanced Fronthaul

As the industry looks toward 6G, the complexity of fronthaul networks will continue to grow. Emerging technologies such as Reconfigurable Intelligent Surfaces (RIS) will introduce new variables into the transport equation, requiring even more precise timing and topology identification.³⁹ The integration of the Y1 interface for exposing radio analytics and the further refinement of intent-based management through the Service Management and Orchestration (SMO) layer will provide the tools needed to manage these future networks.³⁵

The shift toward 6G will also likely see the adoption of even higher functional splits or dynamic splits that adapt to the available transport capacity.⁶ In such an environment, the ability to dynamically identify the transport topology and estimate capacity in real-time will transition from a desirable optimization to a fundamental operational requirement.⁴

Conclusion: A Paradigm for Intelligent Fronthaul Optimization

The optimization of O-RAN fronthaul networks represents a sophisticated intersection of

wireless communications, data science, and high-speed networking. This report has detailed a comprehensive solution for the two primary challenges: network topology identification and link capacity estimation.

By leveraging the statistical correlation of packet loss events, the solution provides a robust and automated method for mapping the "invisible" physical topology of the transport network. This allows for the precise grouping of cells sharing common physical links, which is the foundational step for any dimensioning exercise.

The proposed capacity estimation methodology moves beyond simplistic peak-summing to a more nuanced, statistical approach. By modeling the 4-symbol buffer of the leaf switch using a Leaky Bucket simulation and adhering to the 1% permissible loss threshold, the solution identifies the "sweet spot" of link dimensioning. This approach maximizes the efficiency of the transport infrastructure, significantly reducing TCO without compromising the stringent QoS requirements of 5G NR.

Integrated into the O-RAN architecture via the RIC and xApp framework, and empowered by high-performance technologies like DPDK and Gnuplot, this solution provides a practical, scalable, and intelligent path forward for network operators. As the telecommunications industry continues its journey toward open and virtualized networks, the application of such data-driven optimization techniques will be the key differentiator in building cost-effective, high-performance, and resilient mobile infrastructure.

Works cited

1. Zero-Trust Strategies for O-RAN Cellular Networks: Principles, Challenges and Research Directions - arXiv, accessed on January 31, 2026,
<https://arxiv.org/html/2511.18568v1>
2. Open RAN XHaul: A Guide To Next-Generation Network Transport - Asterfusion, accessed on January 31, 2026,
<https://cloudswit.ch/blogs/everything-about-open-ran-and-xhaul/>
3. Understanding O-RAN: Architecture, Interfaces, Algorithms, Security, and Research Challenges - IEEE Xplore, accessed on January 31, 2026,
<https://ieeexplore.ieee.org/ielam/9739/10130694/10024837-aam.pdf>
4. A Review of Open-RAN Intelligence: Opportunities, Challenges, Real-Life Applications and Impacts | Request PDF - ResearchGate, accessed on January 31, 2026,
https://www.researchgate.net/publication/397801270_A_Review_of_Open-RAN_Intelligence_Opportunities_Challenges_Real-Life_Applications_and_Impacts
5. Hackathon Problem Statement Fronthaul network optimization.pdf
6. A survey of longitudinal changes in cellular network architecture: The good, the bad, and the ugly - ResearchGate, accessed on January 31, 2026,
https://www.researchgate.net/publication/363276503_A_survey_of_longitudinal_changes_in_cellular_network_architecture_The_good_the_bad_and_the_ugly
7. Transport Layer and O-RAN Fronthaul Protocol Implementation — o-du-phy

- master documentation, accessed on January 31, 2026,
https://docs.o-ran-sc.org/projects/o-ran-sc-o-du-phy/en/latest/Transport-Layer-and-ORAN-Fronthaul-Protocol-Implementation_fh.html
8. Exploring functional splits in 5G RAN: Tradeoffs and use cases (Reader Forum), accessed on January 31, 2026,
<https://www.rcrwireless.com/20210317/5g/exploring-functional-splits-in-5g-ran-tradeoffs-and-use-cases-reader-forum>
 9. Data-Driven Self-Management of Cellular Radio Access Networks - RIUMA, accessed on January 31, 2026,
https://riuma.uma.es/xmlui/bitstream/handle/10630/26787/TD_GIJON_MARTIN_Carolina.pdf?sequence=1&isAllowed=y
 10. 5g New Radio Dimensioning | PDF | Antenna (Radio) | Orthogonal Frequency Division Multiplexing - Scribd, accessed on January 31, 2026,
<https://www.scribd.com/document/671801940/5g-New-Radio-Dimensioning>
 11. Development of a C-RAN Fronthaul Simulator - kth .diva, accessed on January 31, 2026, <https://kth.diva-portal.org/smash/get/diva2:1798086/FULLTEXT01.pdf>
 12. CyberPulse: A Security Framework for Software- Defined Networks - VU Research Repository, accessed on January 31, 2026,
https://vuir.vu.edu.au/42172/6/RASOOL_Raihan-Thesis_nosignature.pdf
 13. (PDF) Economic tradeoff in the optimization of carrier aggregation with enhanced multi-band scheduling in LTE-Advanced scenarios - ResearchGate, accessed on January 31, 2026,
https://www.researchgate.net/publication/282560434_Economic_tradeoff_in_the_optimization_of_carrier_aggregation_with_enhanced_multiband_scheduling_in_LTEAdvanced_scenarios
 14. Scheduling allocation in 5G slicing networks utilizing weighted exponential and logarithmic functions to improve QoS - PMC, accessed on January 31, 2026,
<https://pmc.ncbi.nlm.nih.gov/articles/PMC12402267/>
 15. Analyze CU-Plane Messages for O-RAN Fronthaul Testing - MATLAB & Simulink, accessed on January 31, 2026,
<https://www.mathworks.com/help/5g/ug/analyze-cu-plane-messages-for-o-ran-fronthaul-testing.html>
 16. Treatment-Based Classification in Residential Wireless Access Points - Department of Computer Science - Worcester Polytechnic Institute, accessed on January 31, 2026, <http://web.cs.wpi.edu/~claypool/phd/lif/phd.pdf>
 17. A Survey of Graph-Based Resource Management in Wireless, accessed on January 31, 2026, <https://www.ece.uvic.ca/~cai/tccn-survey-part2-2024.pdf>
 18. December 2025 HAROLD VINCENT POOR CONTACT INFORMATION: Electrical and Computer Engineering Phone: +1 (609) 258-1816 Princeton Un, accessed on January 31, 2026, <https://ece.princeton.edu/document/9816>
 19. Distributed Beamforming in MISO SWIPT System | Request PDF - ResearchGate, accessed on January 31, 2026,
https://www.researchgate.net/publication/309148307_Distributed_Beamforming_in_MISO_SWIPT_System
 20. 661 kB - Hugging Face, accessed on January 31, 2026,

<https://huggingface.co/yuningm/bart-large-citesum/resolve/main/pred-3.txt?download=true>

21. On the design of efficient caching systems - UCL Discovery, accessed on January 31, 2026, <https://discovery.ucl.ac.uk/1473436/1/thesis-final.pdf>
22. REAL-TIME O-RAN FRONTHAUL ANALYZER AND GRAPHICAL USER INTERFACE - European Publication Server web service, accessed on January 31, 2026, <https://data.epo.org/publication-server/rest/v1.2/publication-dates/2025-04-09/patents/EP4535855NWA1/document.pdf>
23. Insulation Monitoring and Diagnosis of Faults in Cross-Bonded Cables Based on the Resistive Current and Sheath Current - NDLI, accessed on January 31, 2026, http://www.ndl.gov.in/re_document/doaj/doaj/4f96f5ecf81c426284c66d1d58364251
24. (PDF) Network Coding Aware User Plane for Mobile Networks - ResearchGate, accessed on January 31, 2026, https://www.researchgate.net/publication/346261180_Network_Coding_Aware_User_Plane_for_Mobile_Networks
25. Machine Learning Applied to Optical Communication Systems - MDPI, accessed on January 31, 2026, https://mdpi-res.com/bookfiles/book/11019/Machine_Learning_Applied_to_Optical_Communication_Systems.pdf?v=1749314738
26. A measurement study on the application-level performance of LTE - ResearchGate, accessed on January 31, 2026, https://www.researchgate.net/publication/271481548_A_measurement_study_on_the_application-level_performance_of_LTE
27. Slice Sandwich: Jagged Slicing Multi-Tier Dynamic Resources for Diversified V2X Services, accessed on January 31, 2026, https://www.researchgate.net/publication/378209713_Slice_Sandwich_Jagged_Slicing_Multi-tier_Dynamic_Resources_for_Diversified_V2X_Services
28. Reinforcement Learning-Based Bandwidth Decision in Optical Access Networks: A Study of Exploration Strategy and Time With Confidence Guarantee - ResearchGate, accessed on January 31, 2026, https://www.researchgate.net/publication/381022512_Reinforcement_Learning-based_Bandwidth_Decision_in_Optical_Access_Networks_A_Study_of_Exploration_Strategy_and_Time_with_Confidence_Guarantee
29. Ethernet Passive Optical Networks | Request PDF - ResearchGate, accessed on January 31, 2026, https://www.researchgate.net/publication/258048459_Ethernet_Passive_Optical_Networks
30. Scheduling allocation in 5G slicing networks utilizing weighted exponential and logarithmic functions to improve QoS - ResearchGate, accessed on January 31, 2026, https://www.researchgate.net/publication/395177900_Scheduling_allocation_in_5G_slicing_networks_utilizing_weighted_exponential_and_logarithmic_functions_to_improve_QoS
31. Distributed Network Control Using Machine Learning - eScholarship.org,

- accessed on January 31, 2026,
<https://escholarship.org/content/qt4sz5h3b4/qt4sz5h3b4.pdf>
32. Network reconfiguration and management in 6G telecommunication networks, accessed on January 31, 2026,
<https://www.e-ce.uth.gr/wp-content/uploads/formidable/59/main.pdf>
33. Attack Vector Evaluation and Adversary Simulation using xApps: Utilizing the O-RAN SC Reference Implementation - DN.Lab TH Köln, accessed on January 31, 2026,
https://www.dn.th-koeln.de/wp-content/uploads/2025/06/ma_dieterich_11135427_xapp_signed.pdf
34. doc/ORAN_FHI7.2_Tutorial.md · develop · oai / openairinterface5G · GitLab - Eurecom, accessed on January 31, 2026,
https://gitlab.eurecom.fr/oai/openairinterface5g/-/blob/develop/doc/ORAN_FHI7.2_Tutorial.md
35. 69 New or Updated O-RAN Technical Documents Released since November 2023, accessed on January 31, 2026,
<https://www.o-ran.org/blog/69-new-or-updated-o-ran-technical-documents-released-since-november-2023>
36. RXT-6200 100G Universal Test Module - TelecomTest Solutions, accessed on January 31, 2026,
https://telecomtest.com.au/wp-content/uploads/2018/03/RXT-6200_D05-00-138_P_A00_LR.pdf
37. Towards 5G Enabled Tactile Robotic Telesurgery - arXiv, accessed on January 31, 2026, <https://arxiv.org/pdf/1803.03586>
38. Applications of Artificial Intelligence, Machine Learning and related techniques for Computer Networking Systems - arXiv, accessed on January 31, 2026, <https://arxiv.org/pdf/2105.15103>
39. Center for Pervasive Communications and Computing - CPCC - UC Irvine, accessed on January 31, 2026, <https://cpcc.uci.edu/events.php>